

Amar Dhillon

NYC-Metropolitan Area | amardhillon844@gmail.com | (XXX) XXX-XXXX | linkedin.com/in/amar-dhillon-917537261 | github.com/AmarDhillon05

EDUCATION

University of Maryland, College Park
B.S. in Computer Science, Minor in Statistics; GPA: 3.5

College Park, MD
Expected May 2028

CERTIFICATIONS

Amazon Web Services (AWS) Certified Developer - Associate
Amazon Web Services

Nov 2024 – Nov 2027
Digital Badge

PROJECTS

- IEX Order Book & Matching Engine** | *C++17, STL, POSIX, Linux perf* Jul 2026 – Aug 2026
- Built a C++17 order book and matching engine and replayed **50.9M** real IEX exchange messages through it without ever resyncing to the exchange's book, agreeing with the exchange on 99,122 of 99,123 fills and 100% of top-of-book prices
 - Sustained **1.26M events/sec** over the full replay and a 1.2 μ s per-message p99 on a separate 2M-event run, keeping each price level as a FIFO queue (`std::list`) inside a `std::map`, with an order-ID hash map for $O(1)$ cancels
 - Implemented B+ tree and LSM-tree backends for the event log and benchmarked them on identical 2M-record workloads, where the B+ tree gave **12.8x** the read throughput and 22x faster range queries at the cost of 27% more disk space
 - Traced an apparent 1.65x regression to the CPU frequency governor, not a code change, using perf counters and a do-nothing control book, then cut run-to-run variance from about 6% to 0.06% by pinning cores and fixing the clock

EXPERIENCE

- Software Development Engineer Intern** Sep 2026 – Dec 2026
Amazon Web Services *Redmond, WA*
- Building capacity planning for Amazon Redshift maintenance like patching: Java logic that sets upper and lower bounds on reserved EC2 capacity to balance availability against cost, validated by backtesting

- Software Engineering Intern** Jun 2026 – Aug 2026
Capital One *Richmond, VA*
- Cut AI workflow generation time **65%** to under 15s by replacing a Claude skill's open-ended search with focused sub-skills that call Python tools for workflow search and validation, guided by CloudWatch and OpenSearch latency logs
 - Because the 270K+ row PostgreSQL workflow table had strict access limits, precomputed popular workflows into a category tree that the Claude skill searches first, with Redis caching full workflows to cut database queries

- Software Engineer (Contract)** Jan 2026 – Present
UMD App Development Contracting (Client: Amazon Leo) *College Park, MD*
- Reduced end-to-end latency **34%** for a pipeline routing 6,000+ scraped research documents to inference by replacing S3 staging hops with Kinesis streams feeding Lambda

- Machine Learning Research Intern** Dec 2025 – Present
UMD Institute for Advanced Computer Studies (UMIACS) *College Park, MD*
- For a PhD-led robotics project aimed at ICLR, built a pipeline that overlays AnyGrasp's candidate grasps on camera images in Open3D for a vision-language model, then tuned prompts so the model picks the right grasp more often

- Co-Founder, Software Engineer** Jun 2025 – Present
TerpLabs *College Park, MD*
- Built Redis caching in front of the Prisma ORM queries for Tortuga, the studio's class scheduler with 20,000+ visits in its first 30 days, to reduce database load and latency

TECHNICAL SKILLS

Languages: C++17, C, Python, Java, Rust, SQL, TypeScript, Bash

Systems & Performance: debugging and profiling with Linux perf (hardware counters, top-down analysis), POSIX, STL, B+ trees, LSM trees, data structures and algorithms, benchmarking

Cloud & Tools: AWS (Lambda, Kinesis, EC2, Fargate, CloudWatch monitoring), AWS CDK, PostgreSQL, Redis, Docker, GitHub Actions (CI/CD), Git, Pytest