

# AMAR DHILLON

NYC Area | [amardhillon844@gmail.com](mailto:amardhillon844@gmail.com) | (XXX) XXX-XXXX | [linkedin.com/in/amar-dhillon-917537261](https://www.linkedin.com/in/amar-dhillon-917537261) | [github.com/AmarDhillon05](https://github.com/AmarDhillon05)

## Education

### University of Maryland, College Park

*B.S. in Computer Science, Minor in Statistics; GPA: 3.5*

May 2028

*College Park, MD*

## Experience

### Amazon Web Services

Sep 2026 – Dec 2026

*Software Development Engineer Intern*

*Redmond, WA*

- Balancing Amazon Redshift availability against cost by writing and backtesting Java policies that tune upper/lower bounds on EC2 capacity reserved for operations like patching

### Capital One

Jun 2026 – Aug 2026

*Software Engineering Intern*

*Richmond, VA*

- Rebuilt a legacy test-data workflow builder in React/TypeScript + Electron, adding a React-Flow editor and Jest-tested validation; **90%** of 30+ surveyed developers preferred it to the old tool
- Precomputed a tree-style search index over popular workflows from a 270K+ row PostgreSQL table (Python/SQLAlchemy, Redis cache) to minimize DB queries under access limits
- Cut LLM workflow generation time **65%** to under 15s by decomposing a Claude skill into sub-skills backed by Python search/validation tools, guided by CloudWatch + OpenSearch latency logs

### UMD App Development Contracting (Client: Amazon Leo)

Jan 2026 – Present

*Software Engineer (Contract)*

*College Park, MD*

- Shipped a space-research tool used by Amazon Leo's public-policy team, now maintained by Leo engineers; deployed on AWS Fargate (Docker) via GitHub Actions and CDK (TypeScript)
- Reduced end-to-end latency **34%** for a pipeline routing 6,000+ scraped research documents to inference by replacing S3 staging hops with Kinesis streams feeding Lambda
- Automated fine-tuning of a Bedrock source-relevancy model on an EventBridge schedule, gated to skip redundant retraining runs, with X-Ray tracing and SNS alerts on failure

### UMD Institute for Advanced Computer Studies (UMIACS)

Dec 2025 – Present

*Machine Learning Research Intern*

*College Park, MD*

- Built a Python pipeline for PhysTwin, a PhD-led digital twin project targeting ICLR: runs AnyGrasp on multi-camera images, renders grasps as Open3D overlays for a VLM; tuned the VLM's prompts
- Shortened research-environment setup time **30%** with uv dependency pinning and Docker; added psutil time/memory profiling and Pytest suites to find bottlenecks

### TerpLabs

Jun 2025 – Present

*Co-Founder, Software Engineer*

*College Park, MD*

- Co-founded a 110+ member campus product studio that has shipped 5 live products, including Tortuga, a class scheduler with **20,000+ visits** in its first 30 days; built its Redis caching over Prisma SQL
- Building a RAG meal-suggestion feature for the official UMD app (GPT-4o mini, ChromaDB + SentenceBERT search); **93%** of generated meals rated satisfactory by internal testers

## Projects

### IEX Order Book & Matching Engine | C++17, STL, POSIX, Linux perf

Jul 2026 – Aug 2026

- Engineered a C++17 order book and matching engine that replayed **50.9M** real IEX DEEP+ messages with no resync, matching 99,122 of 99,123 fills exactly and 100% of top-of-book prices
- Benchmarked B+ tree vs LSM skip-list stores for the simulation's event logs on identical 2M-record workloads; the B+ tree gave **12.8×** read throughput and 22× faster range summaries for 27% more disk
- Sustained **1.26M events/sec** end-to-end over the full replay, averaging 686.7 ns/event in the matching engine

## Technical Skills

**Languages:** Python, Java, C++, C, TypeScript/JavaScript, SQL

**Cloud (AWS):** Lambda, Kinesis, Fargate/ECS, API Gateway, EventBridge, Bedrock, CDK, CloudWatch, X-Ray, SNS

**Web & Data:** React, Electron, Node.js, PostgreSQL, Redis, SQLAlchemy, Prisma, OpenSearch, ChromaDB

**AI/ML & Tools:** PyTorch, RAG, Open3D, Docker, GitHub Actions, Linux, Pytest, Jest, uv, perf

**Certification:** AWS Certified Developer – Associate (Nov 2024 – Nov 2027)