

AMAR DHILLON

NYC Area | amardhillon844@gmail.com | (XXX) XXX-XXXX | linkedin.com/in/amar-dhillon-917537261 | github.com/AmarDhillon05

Education

University of Maryland, College Park

B.S. in Computer Science, Minor in Statistics; GPA: 3.5

Expected May 2028

College Park, MD

Projects

IEX Order Book & Matching Engine | C++17, STL, POSIX, Linux perf

Jul 2026 – Aug 2026

- Built a C++17 order book and matching engine and replayed **50.9M** real IEX exchange messages through it without ever resyncing to the exchange's book, matching 99,122 of 99,123 fills and 100% of top-of-book prices
- Implemented B+ tree and LSM-tree backends for the engine's event log and benchmarked them on identical 2M-record workloads: the B+ tree gave **12.8x** the read throughput and 22x faster range summaries for 27% more disk
- Traced an apparent 1.65x slowdown to the CPU frequency governor using perf counters and a null control book, then cut run-to-run variance from about 6% to 0.06% with core pinning and a fixed clock
- Stored price levels in a std::map of std::list FIFOs with a hash map from order ID to list iterator for O(1) cancels, sustaining **1.26M events/sec** over the full replay and a 1.2 μs per-update p99 on a separate 2M-event run

Experience

Amazon Web Services

Sep 2026 – Dec 2026

Software Development Engineer Intern

Redmond, VA

- Writing Java logic that decides how much EC2 capacity Amazon Redshift reserves for operations like patching, setting upper and lower bounds that keep availability high without overspending, and backtesting the policies

Capital One

Jun 2026 – Aug 2026

Software Engineering Intern

Richmond, VA

- Cut AI workflow generation time **65%** to under 15s by splitting one large Claude skill into smaller sub-skills that call Python tools for search and validation, after CloudWatch and OpenSearch latency logs showed where time went
- Because the 270K+ row PostgreSQL workflow table had strict access limits, precomputed popular workflows into a category tree that the Claude skill searches first, with Redis caching full workflows to cut database queries

UMD App Development Contracting (Client: Amazon Leo)

Jan 2026 – Present

Software Engineer (Contract)

College Park, MD

- Reduced end-to-end latency **34%** for a pipeline routing 6,000+ scraped research documents to inference by replacing S3 staging hops with Kinesis streams feeding Lambda

UMD Institute for Advanced Computer Studies (UMIACS)

Dec 2025 – Present

Machine Learning Research Intern

College Park, MD

- For a PhD-led robotics project aimed at ICLR, built a Python pipeline that renders AnyGrasp's candidate grasps on multi-camera images as Open3D overlays for a vision-language model, and tuned its prompts so it picks right more often

TerpLabs

Jun 2025 – Present

Co-Founder, Software Engineer

College Park, MD

- Co-founded a 110+ member student product studio that has shipped 5 live products, and personally built Redis caching in front of the Prisma ORM queries for one, Tortuga, a class scheduler with **20,000+ visits** in its first 30 days

Technical Skills

Languages: C++17, C, Python, Java, Rust, SQL, TypeScript, Bash

Systems & Performance: low-latency C++, debugging and profiling with Linux perf (hardware counters, top-down analysis), POSIX, STL, B+ trees, LSM trees, data structures and algorithms, benchmarking

Cloud & Tools: AWS (Lambda, Kinesis, EC2, Fargate, CloudWatch monitoring), AWS CDK, PostgreSQL, Redis, Docker, GitHub Actions (CI/CD), Git, Pytest

Certification: AWS Certified Developer – Associate (Nov 2024 – Nov 2027)